

CSCI-UA.0480-003
Parallel Computing
Lab# 2

In this second lab you will write parallel code, on CIMS machines, to generate prime numbers from 2 to N inclusive and test scalability and performance.

General notes:

- Program will be written in OpenMP.
- The name of the source code file is: `genprime.c`
- You compile it with: *`gcc -g -Wall -fopenmp -std=c99 -o genprime genprimes.c`*
- Write your program in such a way that the command line is:
`./genprime N t`
Where:
N is a positive number bigger than 2 and less than or equal to 100,000,000
t is the number of threads and is a positive integer that does not exceed 100.
- We will not test your program with N larger than 100,000,000.
- The output of your program is a text file `N.txt` (N is the number entered as argument to your program). That file contains all the prime numbers between 2 and N inclusive, one number per line.
- Assume we will not do any tricks with the input (i.e. We will not deliberately test your program with wrong values of N or t).

The algorithm for generating prime numbers:

There are many algorithms for generating prime numbers and for primality testing. Some are more efficient than others. **For this lab, we will implement the following algorithm**, given N:

1. Generate all numbers from 2 to N.
2. First number is 2, so remove all numbers that are multiple of 2 (i.e. 4, 6, 8, ...). Do not remove the 2 itself.
3. Following number is 3, so remove all multiple of 3 that have not been removed from the previous step. That will be: 9, 15, ... till you reach N.
4. The next number that has not been crossed so far is 5. So, remove all multiple of 5 that have not been crossed before, till you reach N.
5. Continue like this till $\text{floor}((N+1)/2)$.
6. The remaining numbers are the prime numbers.

Note: If you run out of memory in step 1, you can generate numbers from 2 to x, then apply the algorithm to that range, write the solution in the file, then generate numbers from x+1, etc.

Example:

Suppose $N = 20$

floor of $(20+1)/2 = 10 \leftarrow$ where we stop.

Initially we have:

2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20

Let's cross all multiple of 2 (but leave 2):

2, 3, 4, 5, 6, 7, 8, 9, ~~10~~, 11, ~~12~~, 13, ~~14~~, 15, ~~16~~, 17, ~~18~~, 19, ~~20~~

Next number is 3, so we cross all multiple of 3 that have not been crossed:

2, 3, 4, 5, 6, 7, 8, 9, ~~10~~, 11, ~~12~~, 13, ~~14~~, ~~15~~, ~~16~~, 17, ~~18~~, 19, ~~20~~

Next number, larger than 3, that has not been crossed is 5, so we will cross multiple of 5 (i.e. 10, 15, and 20). As you see below, they are all already crossed.

2, 3, 4, 5, 6, 7, 8, 9, ~~10~~, 11, ~~12~~, 13, ~~14~~, ~~15~~, ~~16~~, 17, ~~18~~, 19, ~~20~~

Next number that has not been crossed is 7, so we will cross multiple of 7 (i.e. 14). As you see below, they are all already crossed.

2, 3, 4, 5, 6, 7, 8, 9, ~~10~~, 11, ~~12~~, 13, ~~14~~, ~~15~~, ~~16~~, 17, ~~18~~, 19, ~~20~~

The next number that has not been crossed is 11. This is bigger than 10, so we stop here.

The numbers that have not been crossed are the prime numbers:

2, 3, 5, 7, 11, 13, 17, 19

The file that your program generates is 20.txt and looks like:

```
2
3
5
7
11
13
17
19
```

How to measure the execution time?

Your code will be doing three major parts:

1. Read the input from the command line
2. Generate the prime numbers (as indicated by the algorithm above)
3. Write the output file and exit.

Performance wise, we care only about part 2. Therefore, you do the following:

```
double tstart = 0.0, tend=0.0, ttaken;
```

Part of your program that reads the input from the command line

```
tstart = omp_get_wtime();
```

Part of your program that generates the prime numbers (as indicated by the algorithm above)

```
ttaken = omp_get_wtime() - t_start;
```

```
printf("Time take for the main part: %f\n", ttaken);
```

Write the output file and exit.

Note: Leave the printf in your code. So when we execute your program, the txt file that contains the numbers is generated and the time taken is printed on the screen. **The printf is executed by the master thread only.**

The report

Speedup will be calculated using the time measurement indicated above.

- After you implement the OpenMP version of the above algorithm, generate the following tables:
 - Table 1 (N = 10,000) contains the speedup relative to the running time with 1 thread; the table will have number of threads = 2, 5, 10, 20, and 100.
 - Table 2 (N = 100,000) contains the speedup relative to the running time with 1 thread; the table will have number of threads = 2, 5, 10, 20, and 100.
 - Table 3 (N = 1000,000) contains the speedup relative to the running time with 1 thread; the table will have number of threads = 2, 5, 10, 20, and 100.
 - Table 4 (N = 10,000,000) contains the speedup relative to the running time with 1 thread; the table will have number of threads = 2, 5, 10, 20, and 100.
 - Table 5 (N = 100,000,000) contains the speedup relative to the running time with 1 thread; the table will have number of threads = 2, 5, 10, 20, and 100.
- For each table, explain the behavior that you see. Explain, WHY we see the behavior shown in the table, NOT what the table shows.

What do you have to submit:

A single zip file. The file name is your lastname.firstname.zip

Inside that zip file you need to have:

- genprimes.c
- pdf file containing the tables and explanation.

Submit the file through NYU classes.

What is the grade distribution for this lab?

- Program compiles and run correctly: 5 points.
 - We see some speedups in some problem sizes: 10 points (at least 2x relative to single thread)
 - Tables 25 points
 - Explanation of the tables: 10 points
- Total = 50 points

You will get an automatic (-2) if you do not follow the above steps (file naming, ...).

Very important: After we compile and execute your program, if it does not generate the times (speedup) in your report, you will get a zero in the whole lab for dishonesty. Of course, we will take into account that the times are not the same with every run. But if your program does not compile or crashes and we see that you have numbers in the report, we automatically assume you copied them from somewhere else.

Enjoy!